

The AI Observability Maturity Model

Five levels, from no visibility into what your AI actually said, to a closed loop that catches quality drift before a customer does. **Uptime and correctness are different metrics** — most teams' dashboards only track the first one. Find your level below.

0

/ 4

Blind

AI features run in production with no structured record of what was actually said — only generic application logs that were never built for this.

WHAT'S IN PLACE

Standard app logs: request timing, HTTP status, error stack traces.

THE GAP

No record of the actual input/output pair. A bad answer leaves no trace to investigate.

1

/ 4

Logged, Not Watched

Full input/output pairs are captured somewhere — but nobody, and nothing automated, is actually reviewing them.

WHAT'S IN PLACE

Prompts and responses are logged and retrievable after the fact, if someone thinks to look.

THE GAP

Logging without review is an archive, not observability. Drift is invisible until a user reports it.

2

/ 4

Uptime Confused With Correctness

Real dashboards exist — latency, error rate, cost, throughput. The team believes this counts as AI observability. It doesn't.

WHAT'S IN PLACE

Traditional APM fully wired up: alerts on errors, latency percentiles, uptime SLAs tracked.

THE GAP

A 200 response and a correct response are different things. Nothing here can tell them apart.

3

/ 4

Manual Spot-Checks

A human reviews a sample of outputs on some cadence. It catches real problems — but only the ones that happen to land in the sample.

WHAT'S IN PLACE

Scheduled human review, sampled transcripts, a real (if informal) quality bar being applied.

THE GAP

Doesn't scale with volume, and the time between reviews is a fully unmonitored blind spot.

4

/ 4

Closed-Loop Semantic Monitoring

Every response — or a statistically real sample — is scored automatically for groundedness and safety. Quality drift alerts the same way an error-rate spike would, and a bad output traces back to the exact prompt and model version that produced it.

WHY THIS IS THE TARGET

The gap between "shipped" and "shipped correctly" closes to near zero — and closes automatically, not by luck.